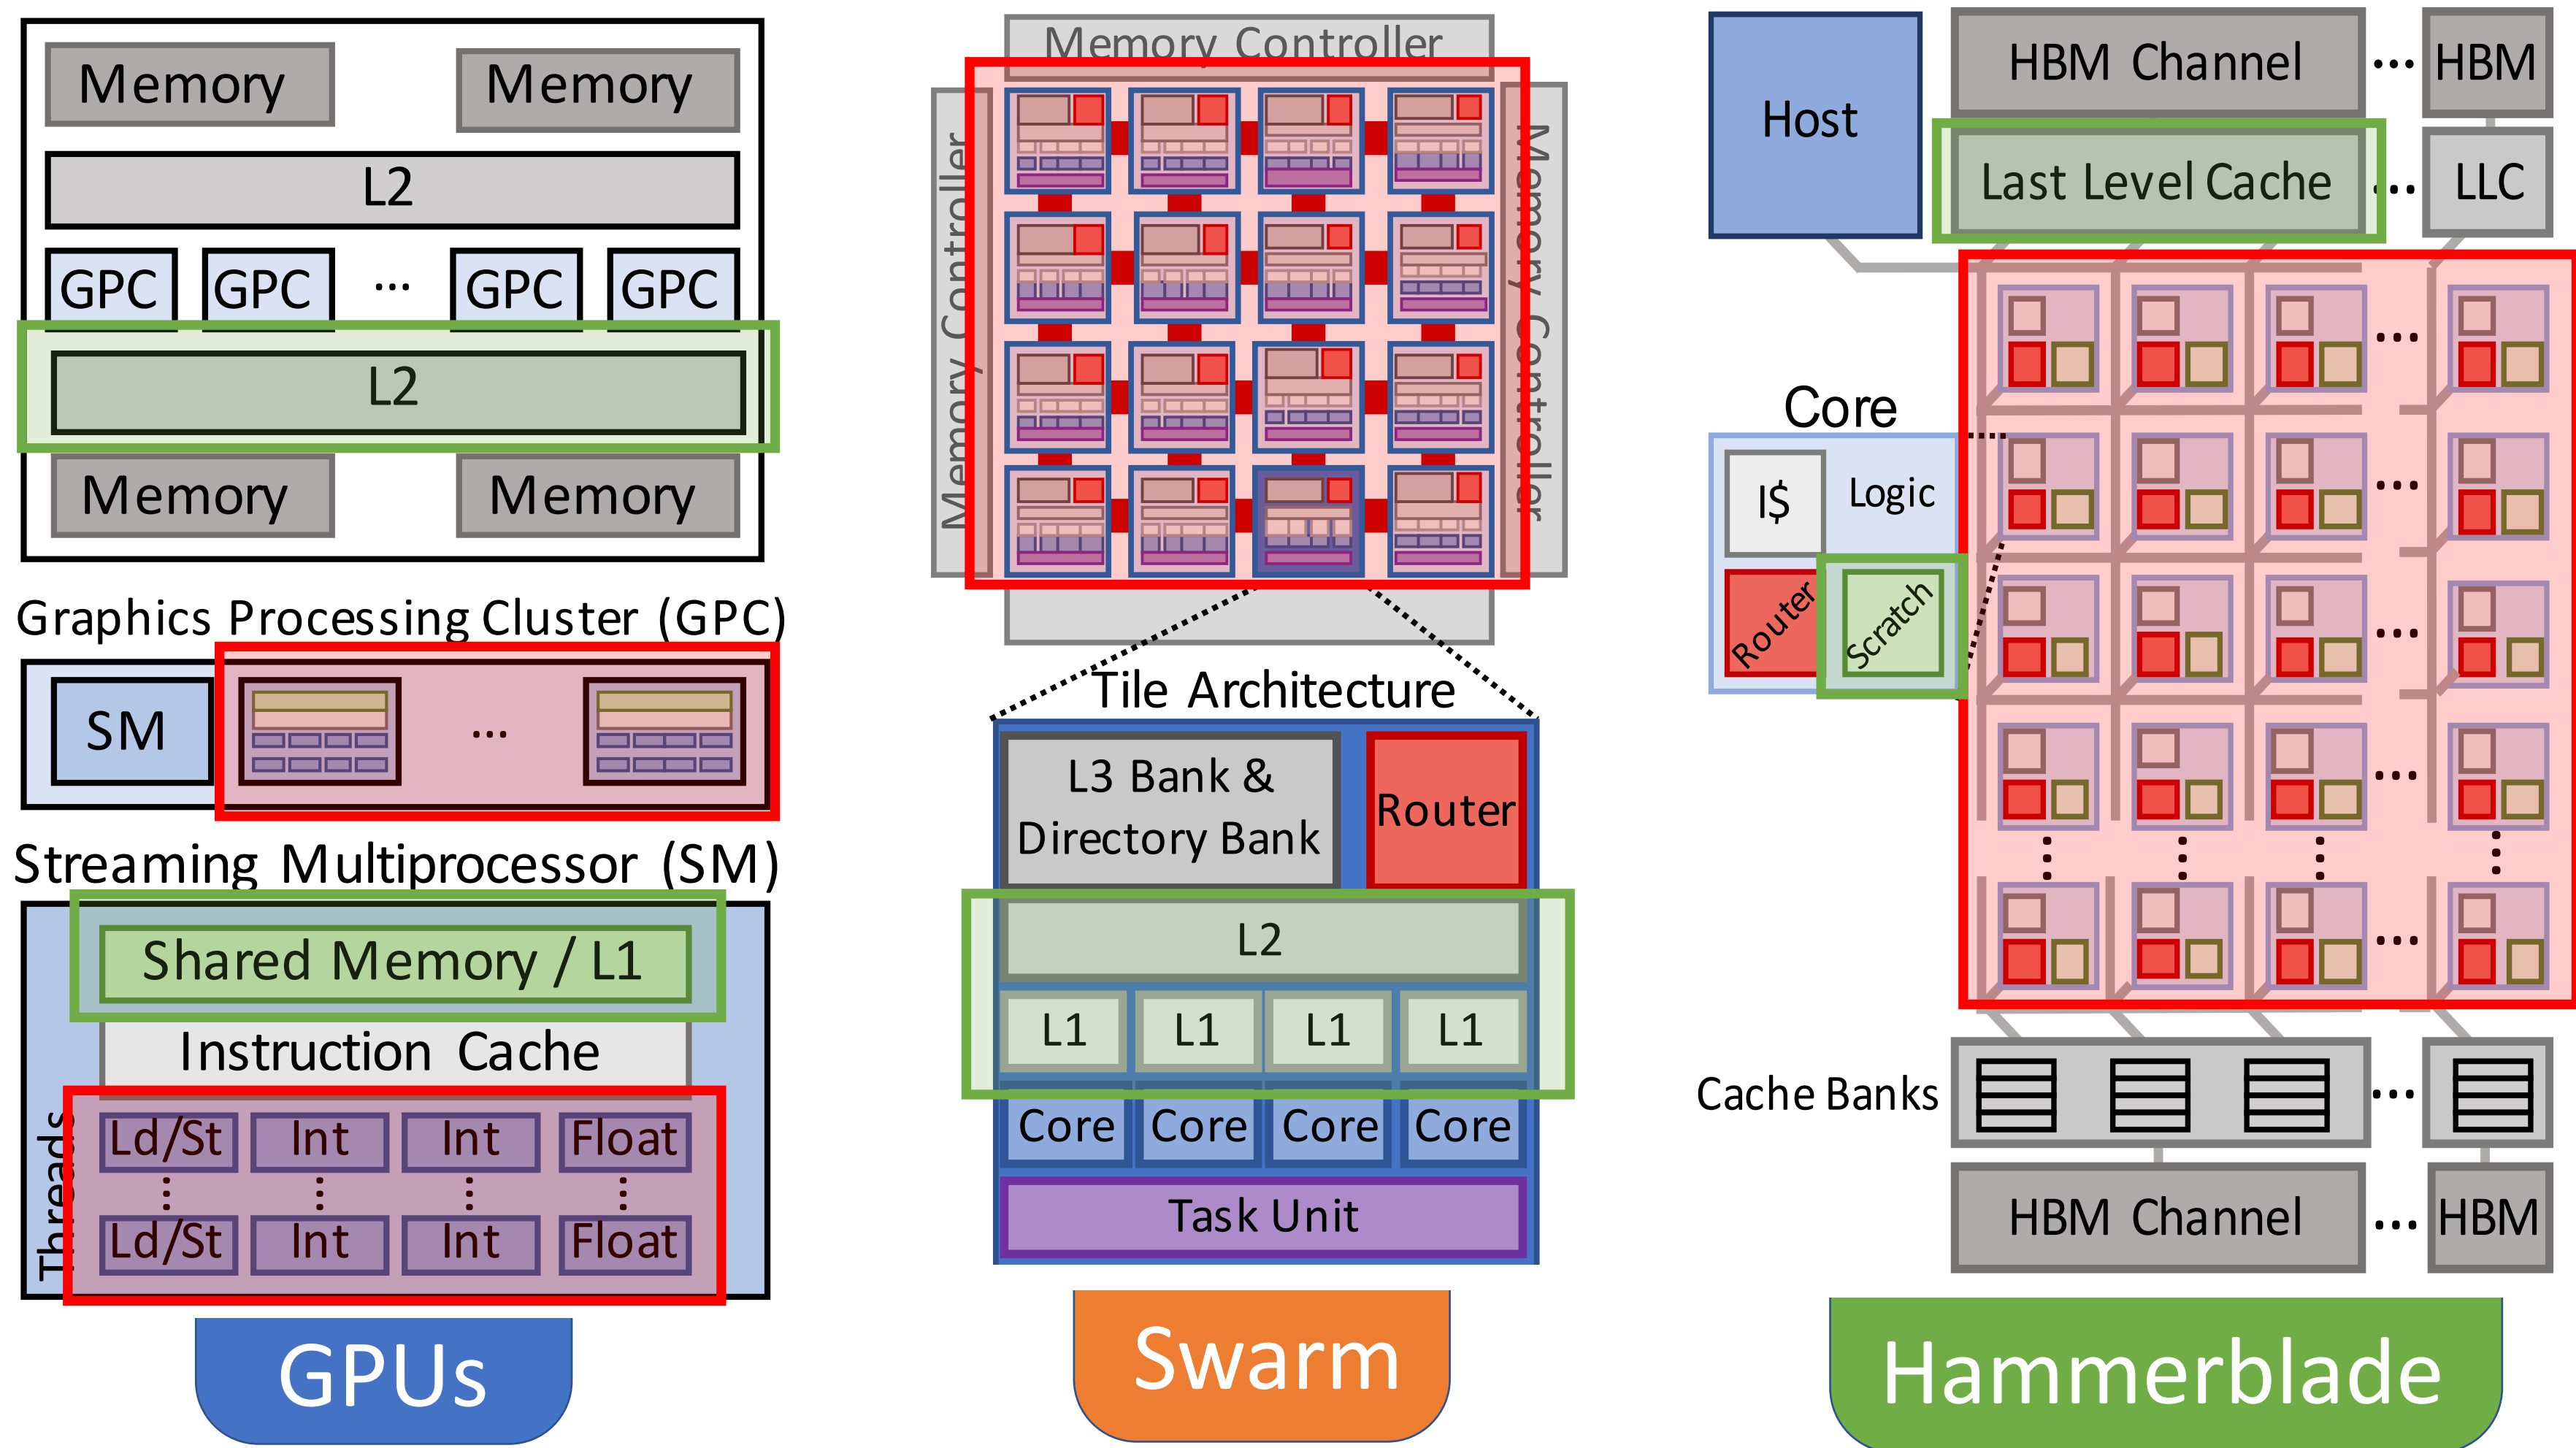


Taming the Zoo: A Unified Graph Compiler Framework for Novel Architectures

Ajay Brahmakshatriya¹, Emily Furst², Victor Yang¹, Claire Hsu¹, Changwan Hong¹, Max Ruttenberg², Yunming Zhang¹, Tommy Jung², Dustin Richmond², Michael Taylor², Julian Shun¹, Mark Oskin², Daniel Sanchez¹, Saman Amarasinghe¹

1 – Massachusetts Institute of Technology, 2 – University of Washington



Similar features in architectures inspire similar optimizations

- Massive parallelism available
 - Load Balance optimizations
- Programmable and mapped caches
 - Blocking for improving memory bandwidth

Application domains require hardware independent optimizations

- PUSH/PULL traversal in graph applications
- Ordered processing for better work efficiency

Single compiler framework for maximizing reusability

```
...
func toFilter(v: Vertex) -> output: bool
    output = (parent[v] == -1);
end
func updateEdge(src: Vertex, dst: Vertex)
    parent[dst] = src;
end
func BFS()
...
#s0# while (frontier.getVertexSetSize() != 0)
    #s1# var output: vertexset{Vertex} =
        edges.from(frontier)
        .to(toFilter)
        .applyModified(updateEdge, parent, true);
...
end
delete frontier
end
GraphIt domain specific program input
Algorithm + Schedule(s)
```

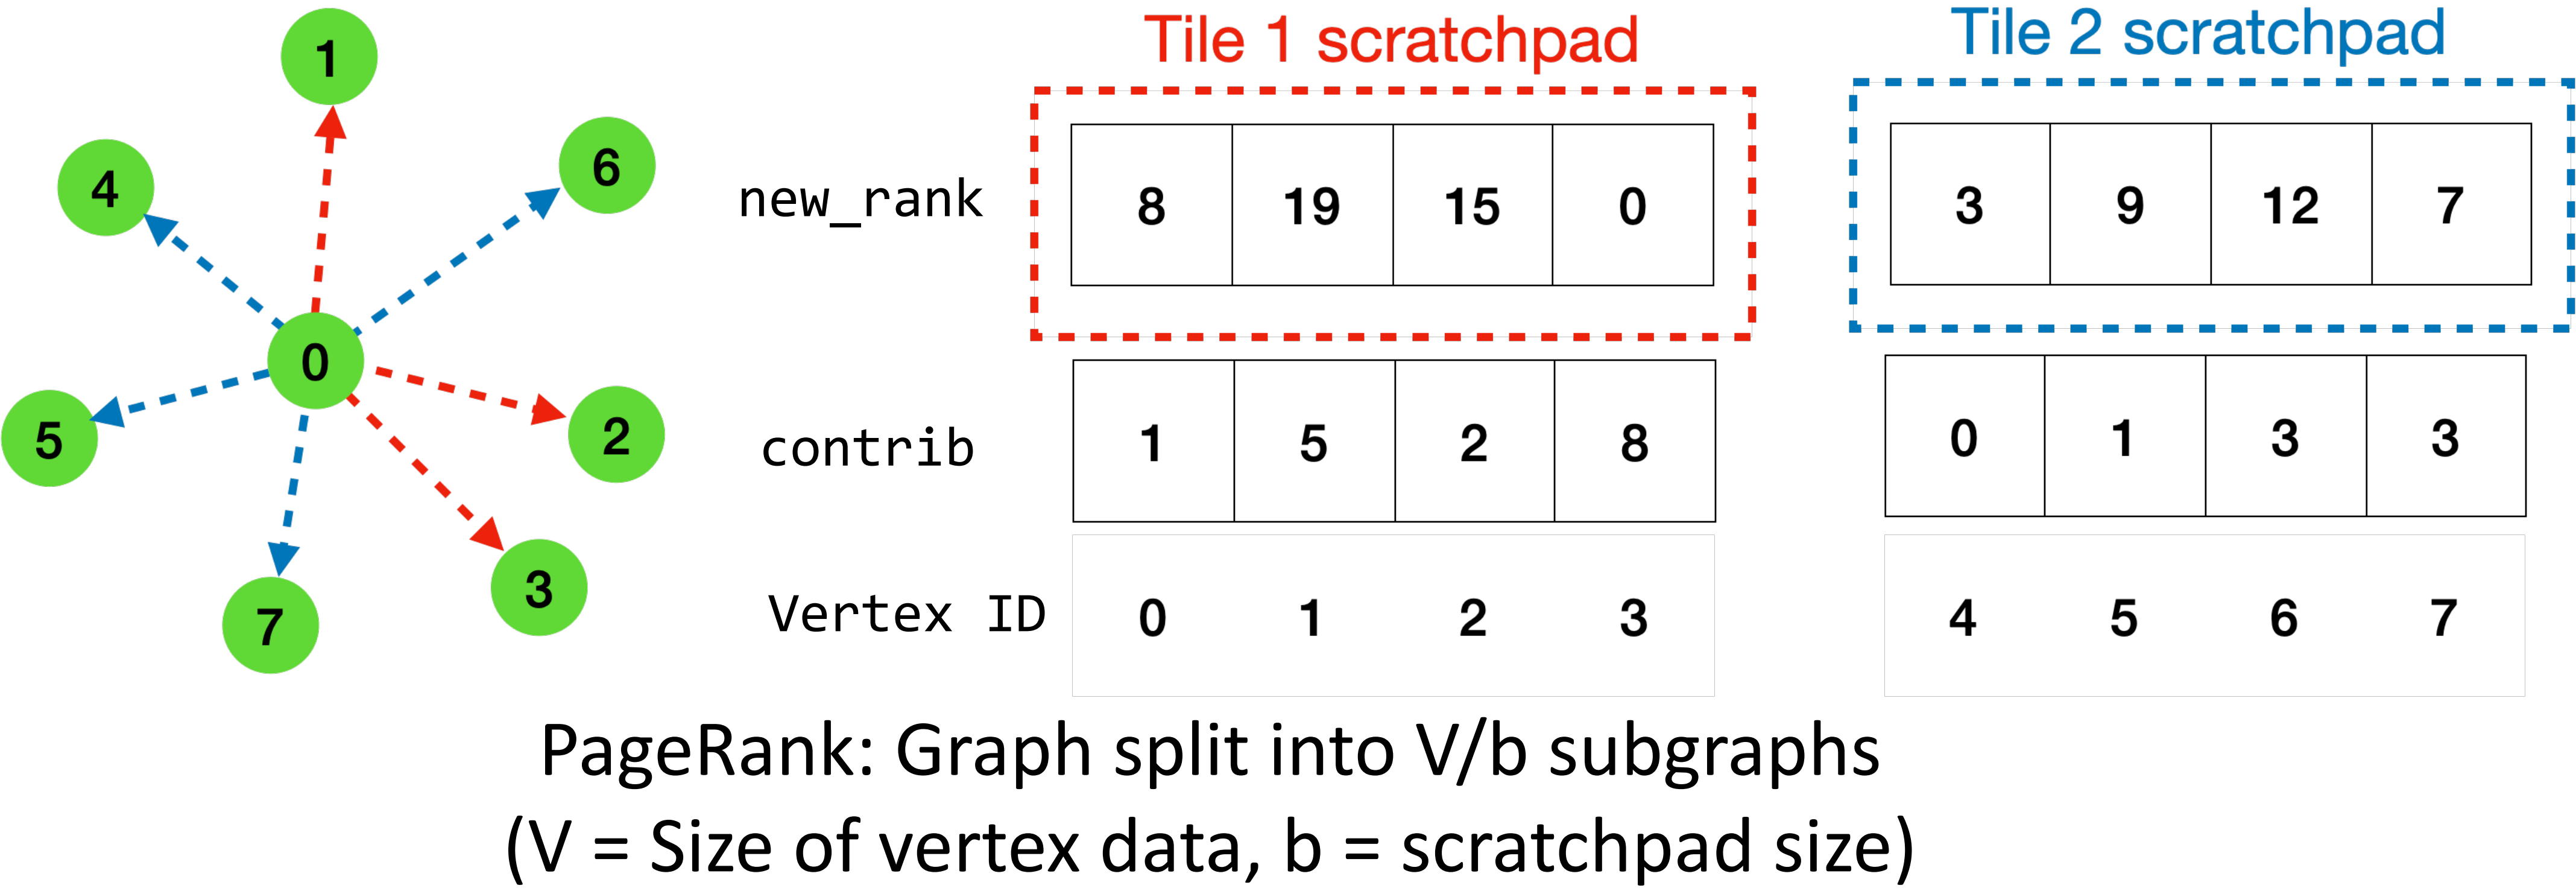
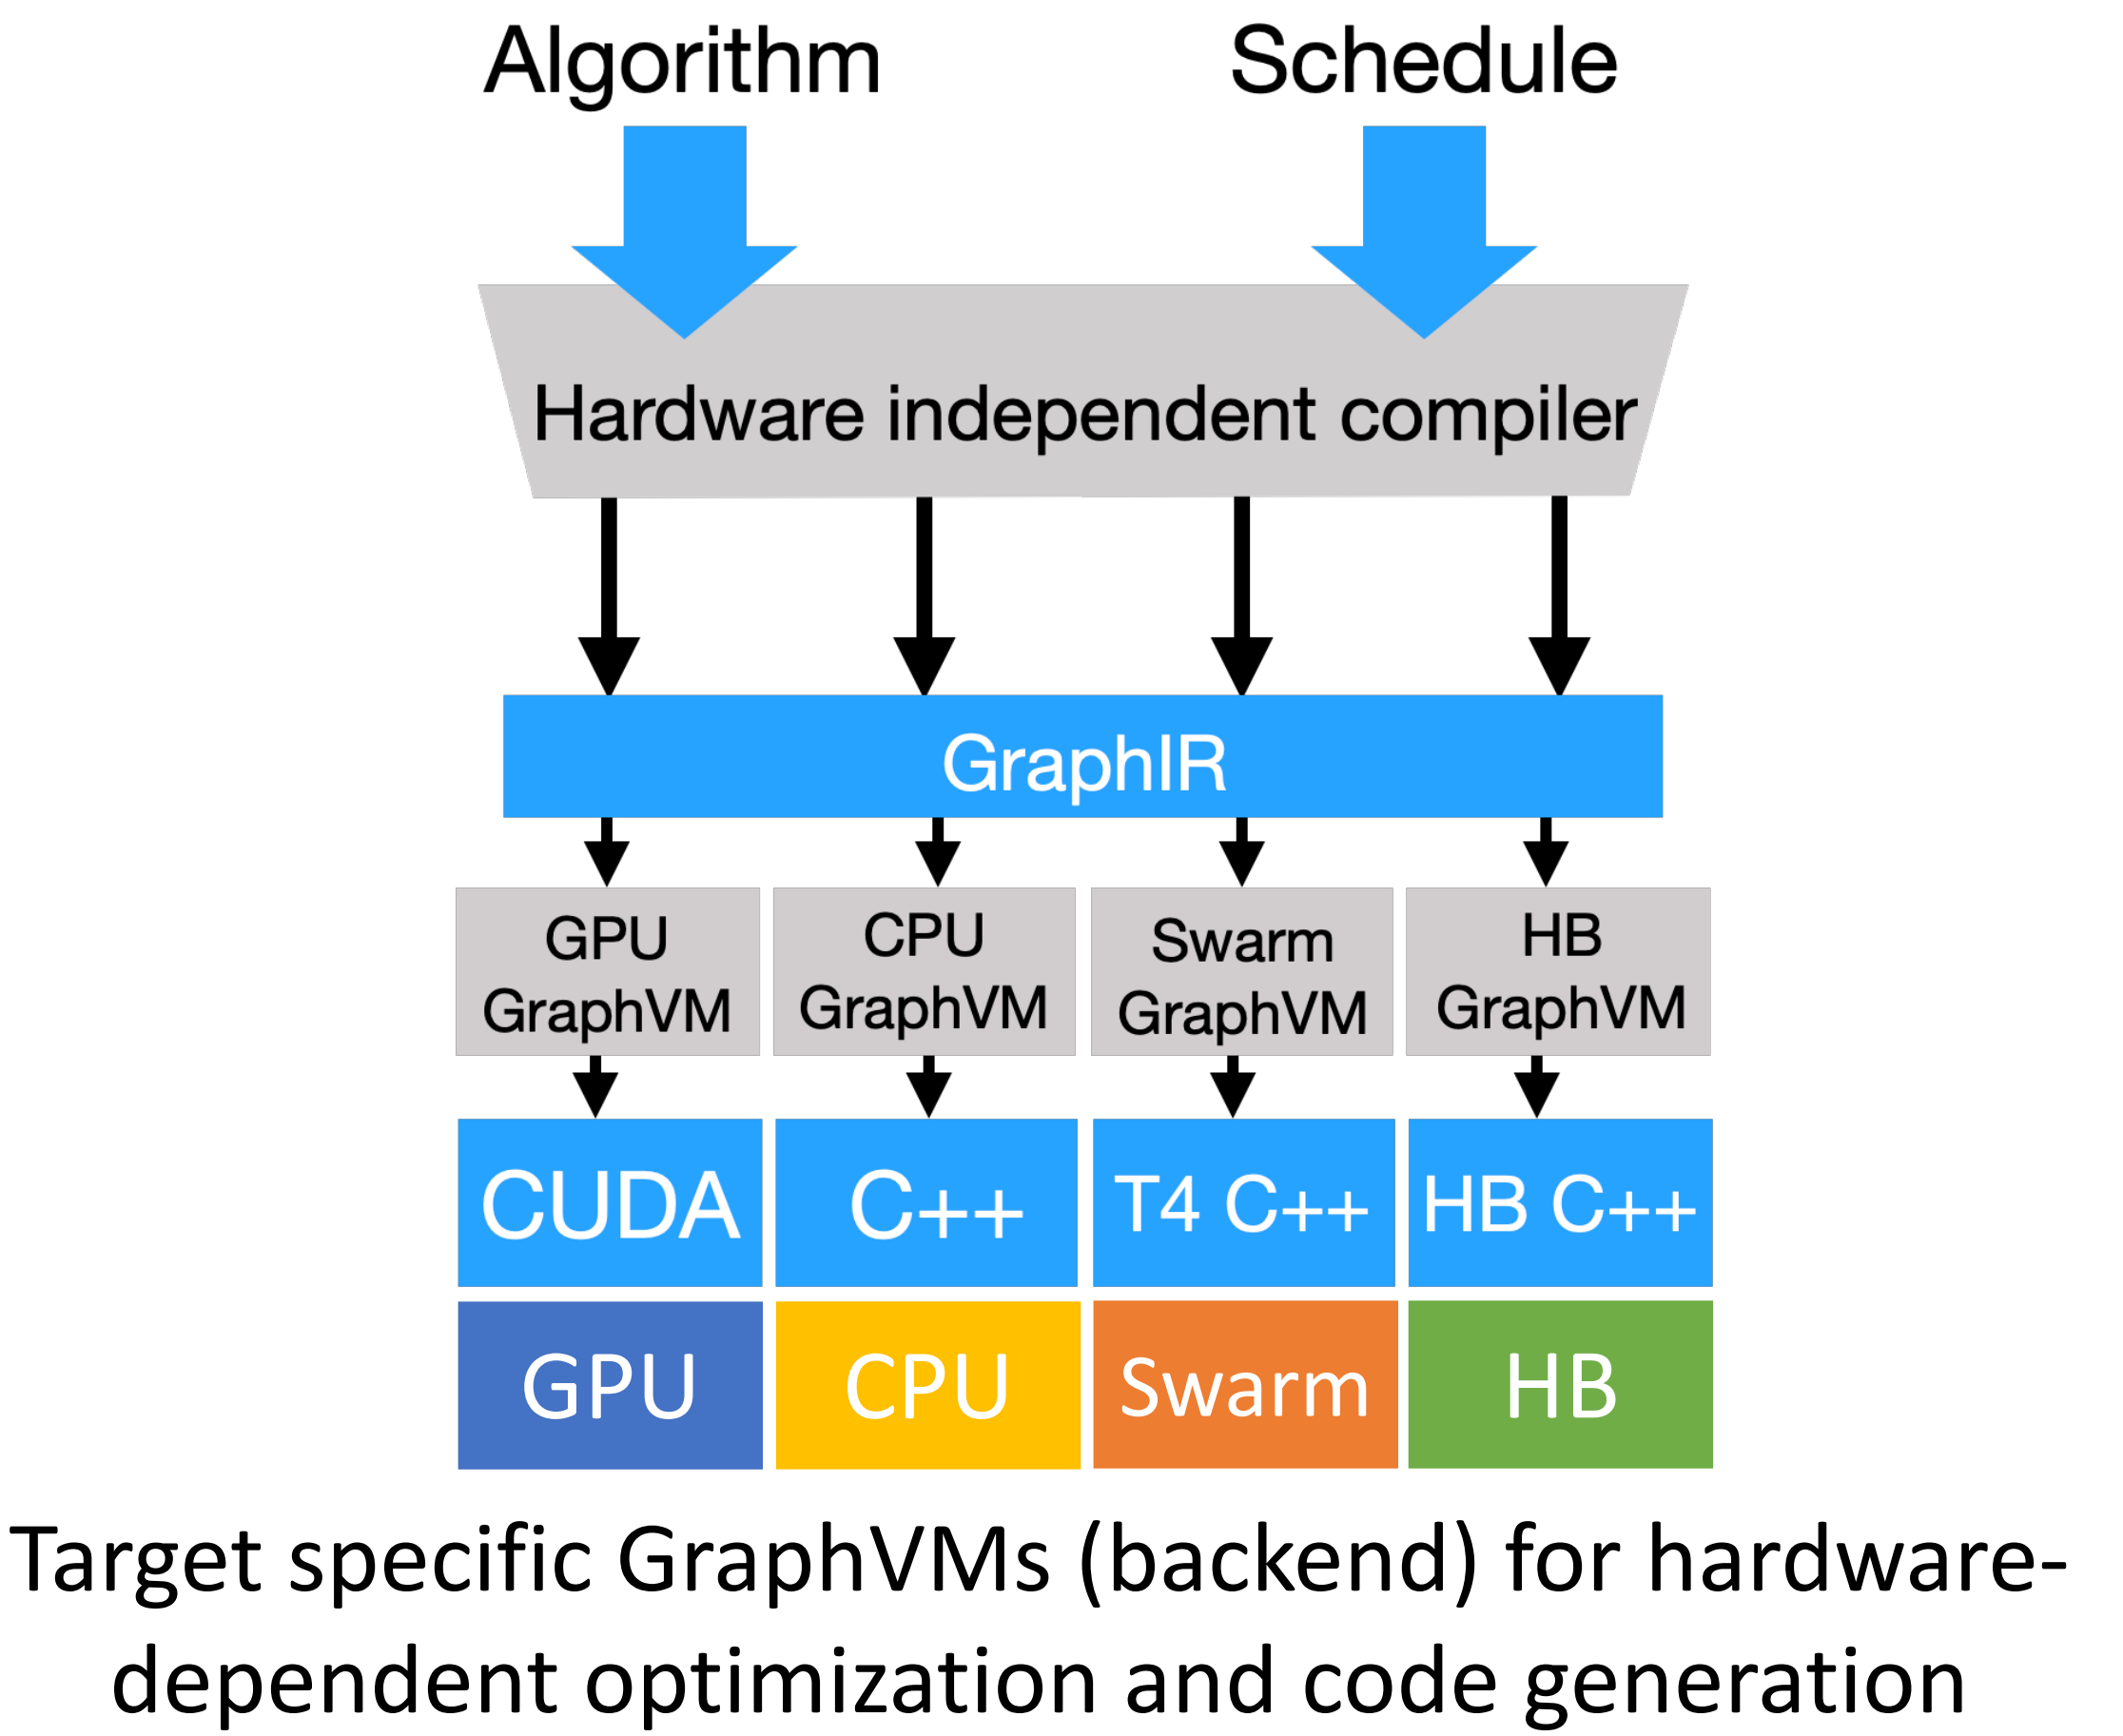
GPU
SimpleGPUSchedule s1;
s1.configDirection(PUSH);
s1.configFrontierCreation(FUSED);
SimpleGPUSchedule s2;
s2.configDirection(PULL, BITMAP);
s2.configFrontierCreation(BITMAP);
CompositeGPUSchedule h1 (INPUT_SET_SIZE, 0.15, s1, s2);

Swarm
SimpleSwarmSchedule s1;
s1.taskGranularity(FINE_GRAINED);
s1.configFrontiers(VERTEXSET_TO_TASKS);

HB
SimpleHBSchedule s1;
s1.configLoadBalance(ALIGNED);
s1.configDirection(HYBRID);

```
Function updateEdge (int32_t src, int32_t dst,
VertexSet output_frontier, {
    bool enqueue = CompareAndSwap<is_atomic=true>(
        parent[dst], -1, src),
    If (enqueue, {
        EnqueueVertex<format=SPARSE>(output_frontier, dst)
    }, {}) }
Function main (int32_t argc, char* argv[], {
    WhileLoopStmt<needs_fusion=true>(VertexSetSize(frontier), {
        EdgeSetIterator<requires_output=true,
            , can_reuse_frontier=true,
            , direction=PUSH,
            , is_edge_parallel=true> (
                edges, frontier, output, updateEdge, toFilter),
        AssignStmt(frontier, output)
    }, )) }
```

Generated Graph Intermediate Representation (GraphIR) with arguments and metadata



Graph blocking optimization implemented for the GPU GraphVM, reused by the Hammerblade GraphVM

Our framework allows for performance critical optimizations with minimum developer effort

| | PR | CC | BC | SSSP | BFS |
|----|------|-------|-------|-------|-------|
| RN | 1.00 | 1.00 | 1.00 | 2.16 | 2.62 |
| RC | 2.11 | 1.00 | 1.00 | 1.97 | 1.45 |
| RU | 1.00 | 1.00 | 1.00 | 2.20 | 3.65 |
| PK | 2.31 | 5.80 | 2.80 | 3.32 | 11.10 |
| HW | 2.34 | 5.23 | 7.09 | 6.86 | 13.41 |
| LJ | 3.04 | 4.09 | 4.40 | 4.60 | 20.14 |
| OK | 3.13 | 3.92 | 5.62 | 6.52 | 52.74 |
| IC | 3.92 | 16.76 | 33.99 | 14.79 | 6.24 |
| TW | 4.43 | 9.53 | 10.66 | 11.59 | 40.61 |
| SW | 3.03 | 2.23 | 2.54 | 3.57 | 5.71 |

GPU (52x)

| | PR | CC | BC | SSSP | BFS |
|----|-------|------|------|------|-------|
| RN | 1.45 | 1.09 | 1.01 | 6.43 | 3.79 |
| RC | 1.53 | 1.15 | 1.00 | 6.89 | 3.38 |
| RU | 1.23 | 1.23 | 1.07 | 3.03 | 3.64 |
| PK | 5.42 | 1.16 | 1.58 | 1.15 | 6.34 |
| HW | 9.59 | 1.27 | 3.43 | 1.00 | 22.31 |
| LJ | 6.74 | 1.27 | 1.80 | 1.16 | 8.62 |
| OK | 7.58 | 1.32 | 4.53 | 1.00 | 42.38 |
| IC | 12.77 | 1.20 | 1.41 | 1.23 | 5.47 |
| TW | 16.05 | 2.44 | 2.40 | 1.00 | 14.32 |
| SW | 38.56 | 6.35 | 8.34 | 1.06 | 42.73 |

CPU (42x)

| | PR | CC | BC | SSSP | BFS |
|----|------|------|------|------|------|
| RN | 5.63 | 8.02 | 1.04 | 1.41 | 1.51 |
| RC | 1.71 | 2.71 | 1.00 | 1.40 | 1.39 |
| RU | 2.78 | 4.95 | 1.00 | 1.43 | 1.42 |
| PK | 2.25 | 1.39 | 1.13 | 1.17 | 1.24 |
| HW | 3.98 | 1.71 | 1.14 | 1.15 | 1.18 |
| LJ | 2.63 | 1.23 | 1.07 | 1.14 | 1.16 |
| OK | 2.41 | 2.41 | 1.15 | 1.16 | 1.36 |
| IC | 6.16 | 1.24 | 1.36 | 1.18 | 1.05 |
| TW | 2.48 | 1.33 | 1.00 | 1.20 | 1.10 |
| SW | 1.33 | 1.72 | 1.00 | 1.02 | 1.06 |

Swarm (8x)

| | PR | CC | BC | SSSP | BFS |
|----|------|------|------|------|------|
| RN | 1.14 | 1.04 | 1.31 | 2.52 | 2.08 |
| RC | 1.02 | 1.00 | 1.07 | 2.18 | 2.32 |
| RU | 1.06 | 1.00 | 1.12 | 2.37 | 2.38 |
| PK | 1.24 | 1.19 | 4.20 | 1.49 | 1.60 |
| HW | 4.97 | | 4.45 | 1.53 | 1.84 |
| LJ | 1.12 | 1.86 | 2.39 | 1.19 | 2.33 |

HB (5x)

| | Base | GPU | CPU | Swarm | HB |
|------------------------|--------|----------|----------|----------|----------|
| Frontend | | | | | |
| Parser + AST | 10,900 | 0 | 0 | 0 | 0 |
| Scheduling Lang. | 136 | +306 | +385 | +524 | +89 |
| Hardware Independent | | | | | |
| Frontier Reuse | 125 | Not Used | 0 | 0 | 0 |
| Property Analysis | 536 | 0 | 0 | Not Used | 0 |
| Ordered Processing | 286 | 0 | +120 | 0 | 0 |
| Other Lowering | 4,171 | 0 | 0 | 0 | 0 |
| Hardware dependent | | | | | |
| Ordered Specialization | 104 | 0 | Not Used | 0 | Not Used |
| Kernel Fusion | 0 | Not Used | +276 | Not Used | Not Used |
| Code Generation | 0 | +3,843 | +1,874 | +959 | +2,282 |
| Runtime Library | 0 | +10,385 | +2,470 | +156 | +1,127 |

LoC for various hardware independent and dependent parts of framework

<https://graphit-lang.org>